

IFT6135 - Practical Assignment One

```
In [3]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import pickle as pickle
%matplotlib inline
```

Data Preparation ¶

Reading data and preparing the test, train and validation sets

```
In [4]: input_size = 784
output_size = 10
```

```

In [5]: class NN:
    ### __init__ Function to load the data and split the train set to train and val
    def __init__(self, mode, datapath=None, hidden_dims=None, hparams={}, model_path=None):
        if mode=='train': assert (hidden_dims is not None) and bool(hparams)
        if mode=='test': assert model_path is not None
        self.mode = mode
        self.model_path = model_path
        self.lr = hparams.get('lr')
        self.weight_decay = hparams.get('weight_decay')
        self.num_epochs = hparams.get('num_epochs')
        self.Params = {}
        data = np.load(datapath)
        # Train data
        self.X_train = data['x_train'][:48000].reshape(48000, 784)
        self.y_train = data['y_train'][:48000]
        # Validation data
        self.X_valid = data['x_train'][48000:].reshape(12000, 784)
        self.y_valid = data['y_train'][48000:]
        self.X_test = data['x_test'].reshape(10000, 784)
        self.y_test = data['y_test']
        if (mode == 'train'):
            self.hidden_dims=(512,512)
            self.init_method = hparams['init_method']
            self.initialize_weights(hidden_dims)
        elif(mode == 'test'):
            with open(model_path,'rb') as f:
                self.Params = pickle.load(f)

    # Initialize the function in three different ways : 1. Normal 2. Glorot 3. Zeros
    def initialize_weights(self, dims):
        n_hidden = len(dims)
        num_nodes = [input_size] + list(dims) + [output_size]
        print('Initialization Method: %s' % self.init_method)
        print('-----\n')
        for i in range(n_hidden + 1):
            self.Params['b%d'%(i+1)] = np.zeros(num_nodes[i+1])
            if self.init_method == 'Normal':
                self.Params['W%d'%(i+1)] = np.random.normal(0, 1,
                                                            (num_nodes[i+1], num_nodes[i]))
            elif self.init_method == 'Glorot':
                D_Glorot = np.sqrt(6.0 / (num_nodes[i+1] + num_nodes[i]))
                self.Params['W%d'%(i+1)] = np.random.uniform(-D_Glorot, D_Glorot,
                                                            (num_nodes[i+1], num_nodes[i]))
            elif self.init_method == 'Zeros':
                self.Params['W%d'%(i+1)] = np.zeros((num_nodes[i+1], num_nodes[i]))
            else:
                raise ValueError('Invalid initialization type: %s' % self.init_method)

    # Forward propagation
    def forward(self, x):
        N = x.shape[0]
        n_hidden = int(len(self.Params.keys()) / 2)
        h_cache_list = []
        a_cache_list = []

```

```

h = x
for i in range(n_hidden):
    h, h_cache = self.fc_forward(h, self.Params['W%d'%(i+1)], self.Params
    h_cache_list.append(h_cache)
    if i != (n_hidden-1):
        h, a_cache = self.relu_forward(h)
        a_cache_list.append(a_cache)
scores = h
cache = (N, n_hidden, a_cache_list, h_cache_list)
return scores, cache

# Backward function
def backward(self, dloss, cache):
    (N, n_hidden, a_cache_list, h_cache_list) = cache
    dx = dloss
    grads = {}
    for i in range(n_hidden, 0, -1):
        dx, dw, db = self.fc_backward(dx, h_cache_list[i-1])
        grads['W%d'%i] = (dw / N) + (self.weight_decay * self.Params['W%d'%i])
        grads['b%d'%i] = (db / N)
        if i != 1:
            dx = self.relu_backward(dx, a_cache_list[i-2])
    return grads

# Computes the forward pass for an affine (fully-connected) layer.
def fc_forward(self, x, w, b):
    out = np.dot(x, w.T) + b
    cache = (x, w, b, out)
    return out, cache

# Computes the backward pass for an affine (fully-connected) layer.
def fc_backward(self, dupstream, cache):
    (x, w, b, out) = cache
    dx = np.dot(dupstream, w)
    dw = np.dot(dupstream.T, x)
    db = dupstream.sum(axis=0)
    return dx, dw, db

### Activation function
def relu_forward(self, x):
    return np.maximum(0, x), x

def relu_backward(self, dupstream, cache):
    return ((cache > 0) * dupstream)

# The function to update parameters
def update(self, grads):
    for param, grad in grads.items():
        self.Params[param] -= self.lr * grad

### Calculate Softmax and Loss using cross entropy
def softmax_loss(self, scores, label):
    sh_log = scores - np.max(scores, axis=1, keepdims=True)
    Z = np.sum(np.exp(sh_log), axis=1, keepdims=True)
    softmax = sh_log - np.log(Z)
    probs = np.exp(softmax)
    dim = scores.shape[0]

```

```

loss = -np.sum(softmax[np.arange(dim), label]) / dim
x = probs.copy()
x[np.arange(dim), label] -= 1
x /= dim
return loss, x

# Train function
def train(self, mb_size=100):
    train_loss_hist = []
    val_loss_hist = []
    train_accuracies = []
    val_accuracies = []
    avg_loss = 0
    for ep in range(self.num_epochs):
        train_loss_epoch = 0
        train_data_size = 0
        correct_pred = 0
        batch_count = np.ceil(self.X_train.shape[0] / mb_size).astype('int')
        for i in range(batch_count):
            xi = self.X_train[i*mb_size:(i+1)*mb_size]
            yi = self.y_train[i*mb_size:(i+1)*mb_size]
            scores, cache = self.forward(xi)
            pred_train = np.argmax(scores, axis=1)
            train_loss, dloss = self.softmax_loss(scores, yi)
            correct_pred += (pred_train == yi).sum()
            grads = self.backward(dloss, cache)
            self.update(grads)
            train_loss_epoch += train_loss * len(yi)
            train_data_size += len(yi)

        scores_val, _ = self.forward(self.X_valid)
        pred_val = np.argmax(scores_val, axis=1)
        val_loss, _ = self.softmax_loss(scores_val, self.y_valid)

        train_loss_hist.append(train_loss_epoch / train_data_size)
        val_loss_hist.append(val_loss)
        train_acc = correct_pred / train_data_size
        val_acc = (pred_val == self.y_valid).mean()
        train_accuracies.append(train_acc)
        val_accuracies.append(val_acc)
        print('Epoch %d' % ep)
        print('Train: Loss: %.4f, Accuracy %.4f' % (train_loss_epoch / train_data_size, train_acc))
        print('Validation: Loss: %.4f, Accuracy %.4f\n' % (val_loss, val_acc))

    train_history = {'train_loss_hist':train_loss_hist,
                    'val_loss_hist':val_loss_hist,
                    'train_accuracies':train_accuracies,
                    'val_accuracies':val_accuracies}
    if self.model_path is not None:
        with open(self.model_path, 'wb') as f:
            pickle.dump(self.Params, f)
    return train_history

# Test function
def test(self, x):
    scores, _ = self.forward(x)
    return scores

```

Helper Functions

```
In [18]: def historyPlt(train_history):
    train_losses = train_history['train_loss_hist']
    train_accuracies = train_history['train_accuracies']
    valid_losses = train_history['val_loss_hist']
    valid_accuracies = train_history['val_accuracies']

    # Plotting the Loss accuracy
    plt.figure(figsize=(12, 4))
    axis = plt.subplot(1, 2, 1)
    axis.plot(range(1, len(train_losses)+1), train_losses, label='train', color =
    axis.plot(range(1, len(valid_losses)+1), valid_losses, label='valid', color =
    axis.legend()
    axis.set_ylabel('Loss')
    axis.set_xlabel('Epochs')

    axis = plt.subplot(1, 2, 2)
    axis.plot(range(1, len(train_accuracies)+1), train_accuracies, label='train',
    axis.plot(range(1, len(valid_accuracies)+1), valid_accuracies, label='valid',
    axis.legend()
    axis.set_ylabel('Accuracy')
    axis.set_xlabel('Epochs')

    def accuracy(scores, labels):
        preds = np.argmax(scores, axis=1)
        accuracy = (preds == labels).mean()
        return accuracy
```

1 .Train the model for 10 epochs 3 using the initialization methods above and record the average loss measured on the training data at the end of each epoch (10 values for each setup).

2. Compare the three setups by plotting the losses against the training time (epoch) and comment on the result.

For this part we trian the model using 3 different initialization approach.

1.1. Initialization using Normal method

```
In [16]: # Initiate hyperparameters
hparams = {'init_method': 'Normal',
           'lr': 1e-1,
           'weight_decay': 0,
           'num_epochs': 10}

# Initialize NN object
nn = NN(mode='train', datapath='./mnist.npz', hidden_dims=(512, 512), hparams=hpa
        model_path='Params.pkl')

# Start training
train_history_normal = nn.train()
print('Training done!')
```

Initialization Method: Normal

Epoch 0

Train: Loss: 766436245291446901776873926236796245558427648.0000, Accuracy 0.1107

Validation: Loss: 2647632.1122, Accuracy 0.1061

Epoch 1

Train: Loss: 2.3023, Accuracy 0.1140

Validation: Loss: 2647632.1121, Accuracy 0.1061

Epoch 2

Train: Loss: 2.3022, Accuracy 0.1140

Validation: Loss: 2647632.1121, Accuracy 0.1061

Epoch 3

Train: Loss: 2.3021, Accuracy 0.1140

Validation: Loss: 2647632.1120, Accuracy 0.1061

Epoch 4

Train: Loss: 2.3020, Accuracy 0.1140

Validation: Loss: 2647632.1119, Accuracy 0.1061

Epoch 5

Train: Loss: 2.3019, Accuracy 0.1140

Validation: Loss: 2647632.1119, Accuracy 0.1061

Epoch 6

Train: Loss: 2.3018, Accuracy 0.1140

Validation: Loss: 2647632.1119, Accuracy 0.1061

Epoch 7

Train: Loss: 2.3017, Accuracy 0.1140

Validation: Loss: 2647632.1118, Accuracy 0.1061

Epoch 8

Train: Loss: 2.3017, Accuracy 0.1140

Validation: Loss: 2647632.1118, Accuracy 0.1061

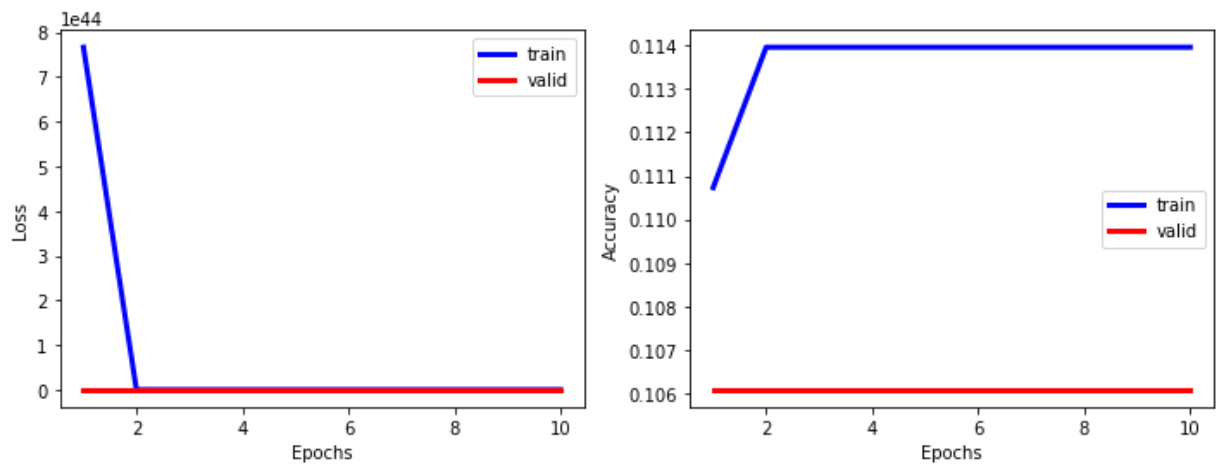
Epoch 9

Train: Loss: 2.3016, Accuracy 0.1140

Validation: Loss: 2647632.1118, Accuracy 0.1061

Training done!

```
In [19]: historyPlt(train_history_normal)
```



1.2. Initialization using Glorot method

```
In [5]: # Initiate hyperparameters
hparams = {'init_method': 'Glorot',
           'lr': 1e-1,
           'weight_decay': 0,
           'num_epochs': 10}

# Initialize NN object
nn = NN(mode='train', datapath='./mnist.npz', hidden_dims=(512, 512), hparams=hpa
        model_path='Params.pkl')

# Start training
train_history_glorot = nn.train()
print('Training done!')
```

Initialization Method: Glorot

Epoch 0

Train: Loss: 3.3376, Accuracy 0.8751

Validation: Loss: 0.8212, Accuracy 0.9147

Epoch 1

Train: Loss: 0.4210, Accuracy 0.9468

Validation: Loss: 0.5856, Accuracy 0.9292

Epoch 2

Train: Loss: 0.2038, Accuracy 0.9662

Validation: Loss: 0.5339, Accuracy 0.9343

Epoch 3

Train: Loss: 0.1037, Accuracy 0.9809

Validation: Loss: 0.4937, Accuracy 0.9378

Epoch 4

Train: Loss: 0.0559, Accuracy 0.9884

Validation: Loss: 0.4795, Accuracy 0.9387

Epoch 5

Train: Loss: 0.0304, Accuracy 0.9940

Validation: Loss: 0.4676, Accuracy 0.9393

Epoch 6

Train: Loss: 0.0163, Accuracy 0.9974

Validation: Loss: 0.4618, Accuracy 0.9413

Epoch 7

Train: Loss: 0.0093, Accuracy 0.9987

Validation: Loss: 0.4565, Accuracy 0.9433

Epoch 8

Train: Loss: 0.0055, Accuracy 0.9995

Validation: Loss: 0.4527, Accuracy 0.9440

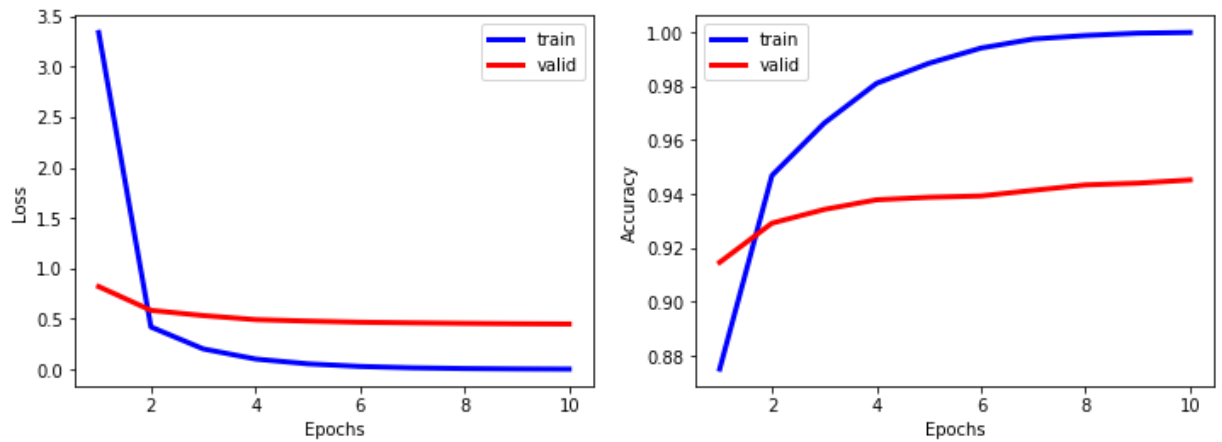
Epoch 9

Train: Loss: 0.0037, Accuracy 0.9998

Validation: Loss: 0.4502, Accuracy 0.9452

Training done!

```
In [52]: historyPlt(train_history_glorot)
```



1.3. Initialization using Zero method

```
In [8]: # Initiate hyperparameters
hparams = {'init_method': 'Zeros',
           'lr': 1e-1,
           'weight_decay': 0,
           'num_epochs': 10}

# Initialize NN object
nn = NN(mode='train', datapath='./mnist.npz', hidden_dims=(512, 512), hparams=hpa
        model_path='Params.pkl')

# Start training
train_history_zero = nn.train()
print('Training done!')
```

Initialization Method: Zeros

Epoch 0

Train: Loss: 2.3025, Accuracy 0.1136

Validation: Loss: 2.3025, Accuracy 0.1060

Epoch 1

Train: Loss: 2.3024, Accuracy 0.1140

Validation: Loss: 2.3024, Accuracy 0.1060

Epoch 2

Train: Loss: 2.3022, Accuracy 0.1140

Validation: Loss: 2.3023, Accuracy 0.1060

Epoch 3

Train: Loss: 2.3021, Accuracy 0.1140

Validation: Loss: 2.3023, Accuracy 0.1060

Epoch 4

Train: Loss: 2.3020, Accuracy 0.1140

Validation: Loss: 2.3022, Accuracy 0.1060

Epoch 5

Train: Loss: 2.3019, Accuracy 0.1140

Validation: Loss: 2.3022, Accuracy 0.1060

Epoch 6

Train: Loss: 2.3018, Accuracy 0.1140

Validation: Loss: 2.3021, Accuracy 0.1060

Epoch 7

Train: Loss: 2.3017, Accuracy 0.1140

Validation: Loss: 2.3021, Accuracy 0.1060

Epoch 8

Train: Loss: 2.3017, Accuracy 0.1140

Validation: Loss: 2.3021, Accuracy 0.1060

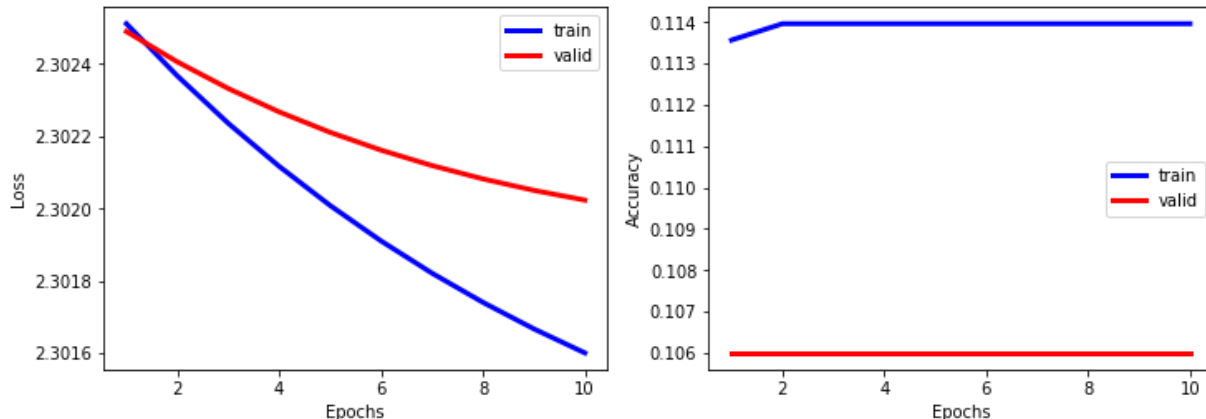
Epoch 9

Train: Loss: 2.3016, Accuracy 0.1140

Validation: Loss: 2.3020, Accuracy 0.1060

Training done!

```
In [53]: historyPlt(train_history_zero)
```



Hyperparameter Search From now on, use the Glorot initialization method.

- 1. Find out a combination of hyper-parameters (model architecture, learning rate, nonlinearity, etc.) such that the average accuracy rate on the validation set ($r(\text{valid})$) is at least 97%.**
- 2. Report the hyper-parameters you tried and the corresponding $r(\text{valid})$.**

For this question we have found two different combination of hyper parameters and initialization methods.

1. using Glorot
2. Using Normal initiation

Optimal Hyperparameter Search

```
In [67]: ## Initiate hyperparameters
# hparams = {'init_method': 'Glorot',
#           'num_epochs': 3}

# for i in range(100):
#     lr = 10 ** np.random.uniform(-2, 0)
#     weight_decay = 10 ** np.random.uniform(-5, 0)

#     hparams['lr'] = lr
#     hparams['weight_decay'] = weight_decay
#     print('lr: %.6f'%lr)
#     print('weight_decay: %.6f'%weight_decay)
#     # Initialize NN object
#     nn = NN((512,512), mode='train', hparams=hparams, datapath='./mnist.npz')

#     # Start training
#     train_history = nn.train()

# print('Training done!')
```

Solution 1 : Hyperparameters Report :

For this question, we trained the model with the following parameters:

- Initialization method : Glorot
- learning rate : 0.14
- batch size : 50
- number of epochs : 10
- weight_decay : 1e-3
- number of neurons for l1, l2 : 600 * 250
- Non linearity : ReLU
- Optimizer : SGD

The case with at least 97% on accuracy rate on the validation set.

```
In [25]: # Initiate hyperparameters
hparams = {'init_method': 'Glorot',
           'lr': 0.14,
           'weight_decay': 1e-3,
           'num_epochs': 10}

# Initialize NN object
nn = NN(mode='train', datapath='./mnist.npz', hidden_dims=(600, 250), hparams=hpa
        model_path='Params.pkl')

# Start training
train_history = nn.train()
print('Training done!')
```

Initialization Method: Glorot

Epoch 0

Train: Loss: 5.6846, Accuracy 0.8661

Validation: Loss: 0.2605, Accuracy 0.9229

Epoch 1

Train: Loss: 0.1918, Accuracy 0.9417

Validation: Loss: 0.1864, Accuracy 0.9453

Epoch 2

Train: Loss: 0.1393, Accuracy 0.9574

Validation: Loss: 0.1550, Accuracy 0.9543

Epoch 3

Train: Loss: 0.1131, Accuracy 0.9663

Validation: Loss: 0.1372, Accuracy 0.9593

Epoch 4

Train: Loss: 0.0964, Accuracy 0.9723

Validation: Loss: 0.1254, Accuracy 0.9617

Epoch 5

Train: Loss: 0.0849, Accuracy 0.9766

Validation: Loss: 0.1175, Accuracy 0.9639

Epoch 6

Train: Loss: 0.0767, Accuracy 0.9798

Validation: Loss: 0.1105, Accuracy 0.9658

Epoch 7

Train: Loss: 0.0714, Accuracy 0.9817

Validation: Loss: 0.1062, Accuracy 0.9678

Epoch 8

Train: Loss: 0.0671, Accuracy 0.9834

Validation: Loss: 0.1028, Accuracy 0.9700

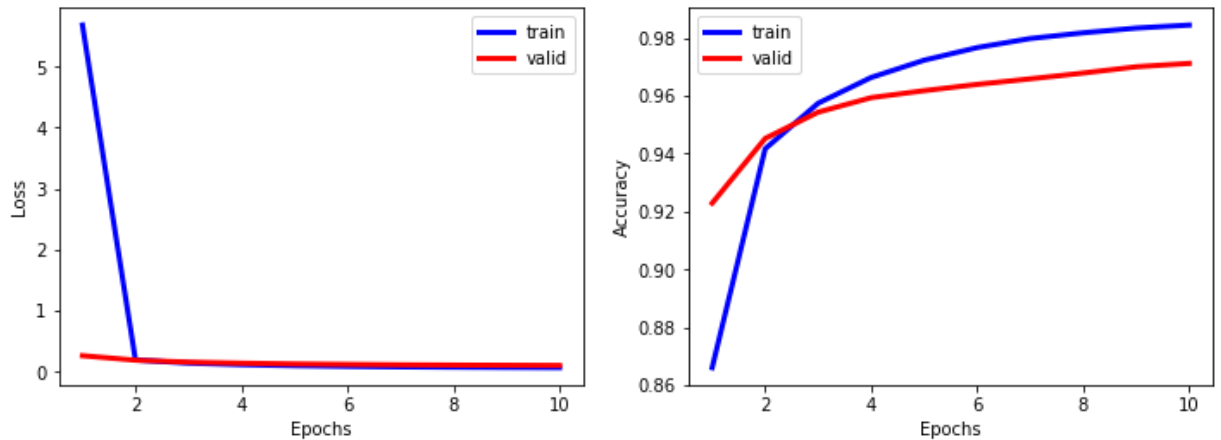
Epoch 9

Train: Loss: 0.0642, Accuracy 0.9844

Validation: Loss: 0.1003, Accuracy 0.9712

Training done!

```
In [54]: historyPlt(train_history)
```



Solution 2 : Hyperparameters Report :

```
In [ ]: For this question, we trained the model with the following parameters:
```

- Initialization method : Normal(mean = 0, sd = 0.001)
- learning rate : 0.14
- batch size : 50
- number of epochs : 10
- weight_decay : 0
- number of neurons for l1, l2 : 512 * 512
- Non linearity : ReLU
- Optimizer : SGD

```
In [33]: # Initiate hyperparameters
hparams = {'init_method': 'Normal',
           'lr': 1e-1,
           'weight_decay': 0,
           'num_epochs': 10}

# Initialize NN object
nn = NN(mode='train', datapath='./mnist.npz', hidden_dims=(512, 512), hparams=hpa
        model_path='Params.pkl')

# Start training
train_history_normal_sd_001 = nn.train()
print('Training done!')
```

Initialization Method: Normal

Epoch 0

Train: Loss: 1.1431, Accuracy 0.6514

Validation: Loss: 0.3999, Accuracy 0.8858

Epoch 1

Train: Loss: 0.2952, Accuracy 0.9145

Validation: Loss: 0.2320, Accuracy 0.9355

Epoch 2

Train: Loss: 0.1824, Accuracy 0.9470

Validation: Loss: 0.1615, Accuracy 0.9554

Epoch 3

Train: Loss: 0.1302, Accuracy 0.9621

Validation: Loss: 0.1291, Accuracy 0.9626

Epoch 4

Train: Loss: 0.0998, Accuracy 0.9718

Validation: Loss: 0.1132, Accuracy 0.9669

Epoch 5

Train: Loss: 0.0793, Accuracy 0.9776

Validation: Loss: 0.1051, Accuracy 0.9690

Epoch 6

Train: Loss: 0.0639, Accuracy 0.9825

Validation: Loss: 0.0996, Accuracy 0.9713

Epoch 7

Train: Loss: 0.0519, Accuracy 0.9861

Validation: Loss: 0.0966, Accuracy 0.9725

Epoch 8

Train: Loss: 0.0424, Accuracy 0.9892

Validation: Loss: 0.0956, Accuracy 0.9726

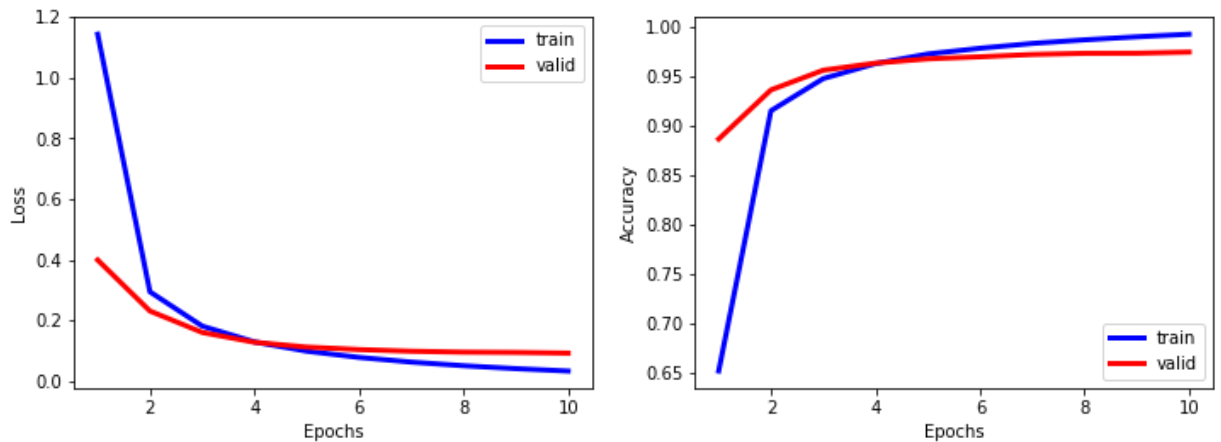
Epoch 9

Train: Loss: 0.0344, Accuracy 0.9918

Validation: Loss: 0.0933, Accuracy 0.9738

Training done!

In [34]: historyPlt(train_history_normal_sd_001)



Numerical Gradient

In [6]:

```
np.random.seed(7899)

hparams = {'init_method': 'Glorot',
           'lr': 1e-1,
           'weight_decay': 2e-3,
           'num_epochs': 10}

nn = NN(mode='train', datapath='./mnist.npz', hidden_dims=(512, 512), hparams=hpa
        model_path='Params.pkl')
xi = nn.X_train[0].reshape(-1,784)
yi = nn.y_train[0].reshape(-1,1)
```

Initialization Method: Glorot

In [13]:

```

max_diffs=[]
Ns=[np.float(i) for i in range(1,25)]
# Ns=[5e0,6e1,7e2,8e3,2e4]
for N in Ns:
    epsi = 1/N
    numerical_grad = []
    for i in range(0,10):

        nn.Params['W2'][0,i] += epsi
        scores,cache = nn.forward(xi)
        softmax_loss_plus, _ = nn.softmax_loss(scores, yi)

        nn.Params['W2'][0,i] -= 2*epsi
        scores,cache = nn.forward(xi)
        softmax_loss_min, _ = nn.softmax_loss(scores, yi)

        numerical_grad.append((softmax_loss_plus - softmax_loss_min) / (2*epsi))

        nn.Params['W2'][0,i] += epsi # reset original value
        scores,cache = nn.forward(xi)
        softmax_loss, dsoftmax = nn.softmax_loss(scores, yi)
        average_grads = nn.backward(dsoftmax, cache)
        diff = np.absolute(average_grads['W2'][0,:10] - numerical_grad)
        max_diffs.append(np.max(diff))

```

In [15]:

```

fig,ax = plt.subplots(1,1)
ax.plot(Ns,max_diffs)
ax.set_ylabel(" Max Diff ")
ax.set_xlabel("N")
plt.show()

```

